──────── Technical Report No. TR-149 ────────

A Tutorial on Spectral Clustering

Ulrike von Luxburg[1]

──────── August 2006 ────────

[1] Department for Empirical Inference, email: ulrike.luxburg@tuebingen.mpg.de

A Tutorial on Spectral Clustering

Ulrike von Luxburg

Abstract. In recent years, spectral clustering has become one of the most popular modern clustering algorithms. It is simple to implement, can be solved efficiently by standard linear algebra software, and very often outperforms traditional clustering algorithms such as the k-means algorithm. Nevertheless, on the first glance spectral clustering looks a bit mysterious, and it is not obvious to see why it works at all and what it really does. This article is a tutorial introduction to spectral clustering. We describe different graph Laplacians and their basic properties, present the most common spectral clustering algorithms, and derive those algorithms from scratch by several different approaches. Advantages and disadvantages of the different spectral clustering algorithms are discussed.

1 Introduction

Clustering is one of the most widely used techniques for exploratory data analysis, with applications ranging from statistics, computer science, biology to social sciences or psychology. In virtually every scientific field dealing with empirical data, people try to get a first impression on their data by trying to identify groups of "similar behavior" in their data. In this article we would like to introduce the reader to the family of spectral clustering algorithms. Compared to the "traditional algorithms" such as k-means or single linkage, spectral clustering has many fundamental advantages. Results obtained by spectral clustering very often outperform the traditional approaches, spectral clustering is very simple to implement and can be solved efficiently by standard linear algebra methods.

This tutorial is set up as a self-contained introduction to spectral clustering. We derive spectral clustering from scratch and present several different points of view to why spectral clustering works. Apart from basic linear algebra, no particular mathematical background is required from the reader. However, we do not attempt to give a concise review of the whole literature on spectral clustering, an endeavor we think would be hopeless due to the overwhelming amount of literature on this subject. The first two sections are devoted to a step-by-step introduction to the mathematical objects used by spectral clustering: similarity graphs in Section 2, and graph Laplacians in Section 3. The spectral clustering algorithms themselves will be presented in Section 4. The next three sections are then devoted to explaining why those algorithms work. Each section corresponds to one explanation: Section 5 describes a graph partitioning approach, Section 6 a random walk perspective, and Section 7 a perturbation theory approach. In Section 8 we will study some practical issues related to spectral clustering, and discuss various extensions and literature related to spectral clustering in Section 9.

2 Graph notation and similarity graphs

Given a set of data points $x_1, \ldots x_n$ and some notion of similarity $s_{ij} \geq 0$ between all pairs of data points x_i and x_j, the intuitive goal of clustering is to divide the data points into several groups such that points in the same group are similar and points in different groups are dissimilar to each other. If we do not have more information than similarities between data points, a nice way of representing the data is in form of the *similarity graph $G = (V, E)$*. The vertices v_i in this graph represent the data points x_i. Two vertices are connected if the similarity s_{ij} between the corresponding data points x_i and x_j is positive (or larger than a certain threshold), and the edge is weighted by s_{ij}. The problem of clustering can now be reformulated using the similarity graph: we want to find a partition of the graph such that the edges between different groups have a very low weight (which means that points in different clusters are dissimilar from each other) and the edges within a group have high weight (which means that points within the same cluster are similar to each other). In this section we want to introduce some basic graph notation, and briefly discuss the kind of graphs we are going to study.

2.1 Graph notation

Let $G = (V, E)$ be an undirected graph with vertex set $V = \{v_1, \dots, v_n\}$. In the following we assume that the graph G is weighted, that is each edge between two vertices v_i and v_j carries a non-negative weight $w_{ij} \geq 0$. The weighted *adjacency matrix* of the graph is the matrix $W = (w_{ij})_{i,j=1,\dots,n}$. If $w_{ij} = 0$ this means that the vertices v_i and v_j are not connected. As G is undirected we require $w_{ij} = w_{ji}$. The degree of a vertex $v_i \in V$ is defined as

$$d_i = \sum_{j=1}^{n} w_{ij}.$$

Note that, in fact, this sum only runs over all vertices adjacent to v_i, as for all other vertices v_j the weight w_{ij} is 0. The *degree matrix* D is defined as the diagonal matrix with the degrees $d_1, \dots, d_n$ on the diagonal. Given a subset of vertices $A \subset V$, we denote its complement $V \setminus A$ by $\overline{A}$. We define the indicator vector $\mathbb{1}_A = (f_1, \dots, f_n)' \in \mathbb{R}^n$ as the vector with entries $f_i = 1$ if $v_i \in A$ and $f_i = 0$ otherwise. For convenience, we introduce the shorthand notation $i \in A$ for the set of indices $\{i|\ v_i \in A\}$. We consider two different ways of measuring the "size" of a subset $A \subset V$:

$$|A| := \text{ the number of vertices in } A$$
$$\text{vol}(A) := \sum_{i \in A} d_i.$$

Intuitively, $|A|$ measures the size of A by its number of vertices, while $\text{vol}(A)$ measures the size of A by the weights of its edges.

A subset $A \subset V$ of a graph is connected if any two vertices in A can be joined by a path such that all intermediate points also lie in A. A subset A is called a connected component if it is connected and if there are no connections between vertices in A and $\overline{A}$. The sets $A_1, \dots, A_k$ form a partition of the graph if $A_i \cap A_j = \emptyset$ and $A_1 \cup \dots \cup A_k = V$.

2.2 Different similarity graphs

There are several popular constructions to transform a given set $x_1, \dots, x_n$ of data points with pairwise similarities s_{ij} or pairwise distances d_{ij} into a graph. The goal when constructing similarity graphs is to model the local neighborhood relationships between the data points. Moreover, most of the constructions below lead to a sparse representation of the data, which has computational advantages.

The ε-neighborhood graph: Here we connect all points whose pairwise distances are smaller than ε. As the distances between all connected points are roughly of the same scale (at most ε), weighting the edges would not incorporate more information about the data to the graph. Hence, the ε-neighborhood graph is usually considered as an unweighted graph.

k-nearest neighbor graphs: Here the goal is to connect vertex v_i with vertex v_j if v_j is among the k nearest neighbors of v_i. However, this definition leads to a directed graph, as the neighborhood relationship is not symmetric. Now there are two ways of making this graph undirected. The first way is to simply ignore the directions of the edges, that is we connect v_i and v_j with an undirected edge if v_i is among the k-nearest neighbors of v_j *or* if v_j is among the k-nearest neighbors of v_i. The resulting graph is what is usually called *the k-nearest neighbor graph*. The second choice is to connect vertices v_i and v_j if both v_i is among the k-nearest neighbors of v_j *and* v_j is among the k-nearest neighbors of v_i. The resulting graph is called the *mutual k-nearest neighbor graph*. In both cases, after connecting the appropriate vertices we weight the edges by the similarity of the adjacent points.

The fully connected graph: Here we simply connect all points with positive similarity with each other, and we weight all edges by s_{ij}. As the graph should model the local neighborhood relationships, this construction is usually only chosen if the similarity function itself already encodes mainly local neighborhoods. An example for a similarity function where this is the case is the Gaussian similarity function $s(x_i, x_j) = \exp(-\frac{\|x_i - x_j\|^2}{2\sigma^2})$. Here the parameter σ controls the width of the neighborhoods, similarly to the parameter ε in case of the ε-neighborhood graph.

All graphs mentioned above are regularly used in spectral clustering. To our knowledge, a thorough study which gives theoretical results about the question which graph one should use under which circumstances does not exist. For an illustration of the behavior of the different graphs we refer to Section 8.

3 Graph Laplacians and their basic properties

The main tools for spectral clustering are graph Laplacian matrices. There exists a whole field dedicated to the study of those matrices, called spectral graph theory (e.g., see Chung, 1997). In this section we want to define different graph Laplacians and point out their most important properties. We will carefully distinguish between different variants of graph Laplacians. Note that in the literature, there is no unique convention which matrix exactly is called "graph Laplacian" and how the different matrices are denoted. Usually, every author just calls "his" matrix the graph Laplacian. Thus, a lot of care is needed when reading literature on graph Laplacians.

In the following we always assume that G is an undirected, weighted graph with weight matrix W, where $w_{ij} = w_{ji} \geq 0$. When we talk about eigenvectors of a matrix, we do not necessarily assume that they are normalized to norm 1. For example, the constant vector $\mathbb{1}$ and a multiple $a\mathbb{1}$ for some $a \neq 0$ are considered as the same eigenvectors. Eigenvalues will always be ordered increasingly, respecting multiplicities. By "the first k eigenvectors" we refer to the eigenvectors corresponding to the k smallest eigenvalues.

3.1 The unnormalized graph Laplacian

The unnormalized graph Laplacian matrix is defined as

$$L = D - W.$$

An overview over many of its properties can be found in Mohar (1991, 1997). The following proposition summarizes the most important facts needed for spectral clustering.

Proposition 1 (Properties of L) *The matrix L satisfies the following properties:*

1. *For every vector $f \in \mathbb{R}^n$ we have*

$$f'Lf = \frac{1}{2} \sum_{i,j=1}^{n} w_{ij}(f_i - f_j)^2.$$

2. *L is symmetric and positive semi-definite.*

3. *The smallest eigenvalue of L is 0, the corresponding eigenvector is the constant one vector $\mathbb{1}$.*

4. *L has n non-negative, real-valued eigenvalues $0 = \lambda_1 \leq \lambda_2 \leq \ldots \leq \lambda_n$.*

Proof.
Part (1): By the definition of d_i,

$$f'Lf = f'Df - f'Wf = \sum_{i=1}^{n} d_i f_i^2 - \sum_{i,j=1}^{n} f_i f_j w_{ij}$$

$$= \frac{1}{2} \left(\sum_{i=1}^{n} d_i f_i^2 - 2 \sum_{i,j=1}^{n} f_i f_j w_{ij} + \sum_{j=1}^{n} d_j f_j^2 \right) = \frac{1}{2} \sum_{i,j=1}^{n} w_{ij}(f_i - f_j)^2.$$

Part (2): The symmetry of L follows directly from the symmetry of W and D. The positive semi-definiteness is a direct consequence of Part (1), which shows that $f'Lf \geq 0$ for all $f \in \mathbb{R}^n$.

Part (3): Obvious.

Part (4) is a direct consequence of Parts (1) - (3).   ☺

Note that the unnormalized graph Laplacian does not depend on the diagonal elements of the adjacency matrix W. Each matrix U which coincides with W on all off-diagonal positions leads to the same unnormalized graph Laplacian L. So in particular, self-edges in a graph do not change the corresponding graph Laplacian.

The unnormalized graph Laplacian and its eigenvalues and eigenvectors can be used to describe many properties of graphs, see Mohar (1991, 1997). One example which will be important for spectral clustering is the following proposition:

Proposition 2 (Number of connected components) *Let G be an undirected graph with non-negative weights. Then the multiplicity k of the eigenvalue 0 of L equals the number of connected components $A_1, \ldots, A_k$ in the graph. The eigenspace of eigenvalue 0 is spanned by the indicator vectors $\mathbb{1}_{A_1}, \ldots, \mathbb{1}_{A_k}$ of those components.*

Proof. We start with the case $k = 1$, that is the graph is connected. Assume that f is an eigenvector with eigenvalue 0. Then we know that

$$0 = f'Lf = \sum_{i,j=1}^{n} w_{ij}(f_i - f_j)^2.$$

As the weights w_{ij} are non-negative, this sum can only vanish if all terms $w_{ij}(f_i - f_j)^2$ vanish. Thus, if two vertices v_i and v_j are connected (i.e., $w_{ij} > 0$), then f_i needs to equal f_j. With this argument we can see that f needs to be constant on the whole connected component. In the case where the graph is connected, we thus only have the constant one vector $\mathbb{1}$ as eigenvector with eigenvalue 0, which obviously is the indicator vector of the connected component.

Now consider the case of k connected components. Without loss of generality we assume that the vertices are ordered according to the connected components they belong to. In this case, the adjacency matrix W has a block diagonal form, and the same is true for the matrix L:

$$L = \begin{pmatrix} L_1 & & & \\ & L_2 & & \\ & & \ddots & \\ & & & L_k \end{pmatrix}$$

Note that each of the blocks L_i is a proper graph Laplacian on its own, namely the Laplacian corresponding to the subgraph of the i-th connected component. As it is the case for all block diagonal matrices, we know that the spectrum of L is given by the union of the spectra of L_i, and the corresponding eigenvectors of L are the eigenvectors of L_i, filled with 0 at the positions of the other blocks. As each L_i is a graph Laplacian of a connected graph, we know that every L_i has eigenvalue 0 with multiplicity 1, and the corresponding eigenvector is the constant one vector on the i-th connected component. Thus, the matrix L has as many eigenvalues 0 as there are connected components, and the corresponding eigenvectors are the indicator vectors of the connected components. ☺

3.2 The normalized graph Laplacians

There are two matrices which are called normalized graph Laplacians in the literature. Both matrices are closely related to each other and are defined as

$$L_{\text{sym}} := D^{-1/2}LD^{-1/2} = I - D^{-1/2}WD^{-1/2}$$
$$L_{\text{rw}} := D^{-1}L = I - D^{-1}W.$$

We denote the first matrix by L_{sym} as it is a symmetric matrix, and the second one by L_{rw} as it is closely connected to a random walk. The standard reference for normalized graph Laplacians is Chung (1997). In the following we summarize several properties of L_{sym} and L_{rw}.

Proposition 3 (Properties of L_{sym} and L_{rw}) *The normalized Laplacians satisfy the following properties:*

1. For every $f \in \mathbb{R}^n$ we have

$$f' L_{sym} f = \frac{1}{2} \sum_{i,j=1}^{n} w_{ij} \left(\frac{f_i}{\sqrt{d_i}} - \frac{f_j}{\sqrt{d_j}} \right)^2 .$$

2. λ is an eigenvalue of L_{rw} with eigenvector v if and only if λ is an eigenvalue of L_{sym} with eigenvector $w = D^{1/2}v$.

3. λ is an eigenvalue of L_{rw} with eigenvector v if and only if λ and v solve the generalized eigenproblem $Lv = \lambda D v$.

4. 0 is an eigenvalue of L_{rw} with the constant one vector $\mathbb{1}$ as eigenvector. 0 is an eigenvalue of L_{sym} with eigenvector $D^{1/2}\mathbb{1}$.

5. L_{sym} and L_{rw} are positive semi-definite and have n non-negative real-valued eigenvalues $0 = \lambda_1 \leq \ldots \leq \lambda_n$.

Proof. Part (1) can be proved similarly to Part (1) of Proposition 1.

Part (2) can be seen immediately by multiplying the eigenvalue equation $L_{\mathrm{sym}} w = \lambda w$ with $D^{-1/2}$ from left and substituting $v = D^{-1/2}w$.

Part (3) follows directly by multiplying the eigenvalue equation $L_{\mathrm{rw}} v = \lambda v$ with D from left.

Part (4) is obvious as $L_{\mathrm{rw}} \mathbb{1} = 0$.

Part (5): The statement about L_{sym} follows from (1), and then the statement about L_{rw} follows from (2).   ☺

As it is the case for the unnormalized graph Laplacian, the multiplicity of the eigenvalue 0 of the normalized graph Laplacian is related to the number of connected components:

Proposition 4 (Number of connected components) *Let G be an undirected graph with non-negative weights. Then the multiplicity k of the eigenvalue 0 of both L_{rw} and L_{sym} equals the number of connected components $A_1, \ldots, A_k$ in the graph. For L_{rw}, the eigenspace of 0 is spanned by the indicator vectors $\mathbb{1}_{A_i}$ of those components. For L_{sym}, the eigenspace of 0 is spanned by the vectors $D^{1/2}\mathbb{1}_{A_i}$.*

Proof. The proof is analogous to the one of Proposition 2, using Proposition 3.   ☺

4   Spectral Clustering Algorithms

Now we would like to state the most common spectral clustering algorithms, for many references and the history of spectral clustering we refer to Section 9. We assume that we are given data points $x_1, \ldots, x_n$ which can be arbitrary objects, and their similarities $s_{ij} = s(x_i, x_j)$, measured according to some similarity function which is symmetric and non-negative. We denote the corresponding similarity matrix by $S = (s_{ij})_{i,j=1\ldots n}$.

Unnormalized spectral clustering

```
Input:  Similarity matrix S ∈ ℝ^{n×n}, number k of clusters to construct
```
- `Construct a similarity graph by one of the ways described in Section 2. Let W be its weighted adjacency matrix.`
- `Compute the unnormalized Laplacian L.`
- **Compute the first k eigenvectors $v_1, \ldots, v_k$ of L.**
- `Let V ∈ ℝ^{n×k} be the matrix containing the vectors v_1,...,v_k as columns.`
- `For i = 1,...,n, let y_i ∈ ℝ^k be the vector corresponding to the i-th row of V.`
- `Cluster the points (y_i)_{i=1,...,n} in ℝ^k with the k-means algorithm into clusters C_1,...,C_k.`
```
Output:  Clusters A_1,...,A_k with A_i = {j| y_j ∈ C_i}.
```

There are two different versions of normalized spectral clustering, depending which of the normalized graph Laplacians is used. We name both algorithms after two popular papers, for more references and history please see Section 9.

Normalized spectral clustering according to Shi and Malik (2000)

Input: Similarity matrix $S \in \mathbb{R}^{n \times n}$, number k of clusters to construct
- Construct a similarity graph by one of the ways described in Section 2. Let W be its weighted adjacency matrix.
- Compute the unnormalized Laplacian L.
- **Compute the first k eigenvectors $v_1, \ldots, v_k$ of the generalized eigenproblem $Lv = \lambda Dv$.**
- Let $V \in \mathbb{R}^{n \times k}$ be the matrix containing the vectors $v_1, \ldots, v_k$ as columns.
- For $i = 1, \ldots, n$, let $y_i \in \mathbb{R}^k$ be the vector corresponding to the i-th row of V.
- Cluster the points $(y_i)_{i=1,\ldots,n}$ in $\mathbb{R}^k$ with the k-means algorithm into clusters $C_1, \ldots, C_k$.

Output: Clusters $A_1, \ldots, A_k$ with $A_i = \{j | y_j \in C_i\}$.

Note that this algorithm uses the generalized eigenvectors of L, which according to Proposition 3 correspond to the eigenvectors of the matrix L_{rw}. So in fact, the algorithm works with eigenvectors of the normalized Laplacian L_{rw}, and hence is called normalized spectral clustering. The next algorithm also uses a normalized Laplacian, but this time the matrix L_{sym} instead of L_{rw}. As we will see, this algorithm needs to introduce an additional row normalization step which is not needed in the the other algorithms. The reasons will become clear in Section 7.

Normalized spectral clustering according to Ng, Jordan, and Weiss (2002)

Input: Similarity matrix $S \in \mathbb{R}^{n \times n}$, number k of clusters to construct
- Construct a similarity graph by one of the ways described in Section 2. Let W be its weighted adjacency matrix.
- Compute the normalized Laplacian L_{sym}.
- **Compute the first k eigenvectors $v_1, \ldots, v_k$ of L_{sym}.**
- Let $V \in \mathbb{R}^{n \times k}$ be the matrix containing the vectors $v_1, \ldots, v_k$ as columns.
- **Form the matrix $U \in \mathbb{R}^{n \times k}$ from V by normalizing the row sums to have norm 1,** that is $u_{ij} = v_{ij}/(\sum_k v_{ik}^2)^{1/2}$.
- For $i = 1, \ldots, n$, let $y_i \in \mathbb{R}^k$ be the vector corresponding to the i-th row of U.
- Cluster the points $(y_i)_{i=1,\ldots,n}$ with the k-means algorithm into clusters $C_1, \ldots, C_k$.

Output: Clusters $A_1, \ldots, A_k$ with $A_i = \{j | y_j \in C_i\}$.

All three algorithm stated above look rather similar, apart from the fact that they use the three different graph Laplacians. In all three algorithms, the main trick is to change the representation of the abstract data points x_i to points $y_i \in \mathbb{R}^k$. It is due to the properties of the graph Laplacians that this change of representation is useful. We will see in the next sections that this change of representation enhances the cluster-properties in the data, so that they can be trivially detected in the new representation. In particular, the simple k-means clustering algorithm has no difficulties to detect the clusters in this new representation. Readers not familiar with k-means can read up on this algorithm for example in Hastie, Tibshirani, and Friedman (2001).
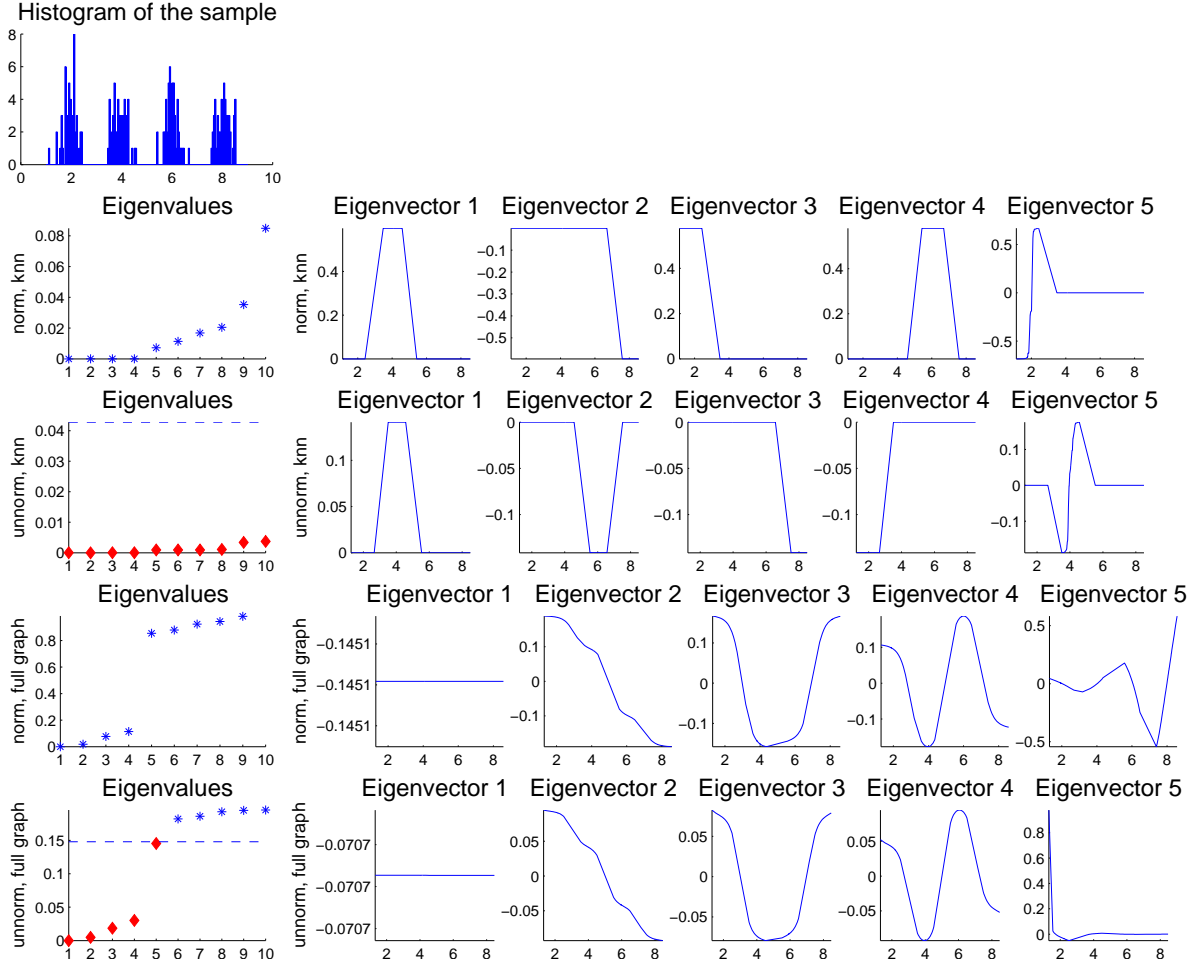
Figure 1: Toy example for spectral clustering. Left upper corner: histogram of the data. First and second row: eigenvalues and eigenvectors of L_{rw} and L based on the k-nearest neighbor graph. Third and fourth row: eigenvalues and eigenvectors of L_{rw} and L based on the fully connected graph. For all plots, we used we use the Gaussian kernel with $\sigma = 1$ as similarity function.

Before we dive into the theory of spectral clustering, we would like to illustrate its principle on a very simple toy example. This example will be used at several places in this tutorial, and we chose it because it is so simple that the relevant quantities can easily be plotted. This toy data set consists of a random sample of 200 points $x_1, \ldots, x_{200} \in \mathbb{R}$ drawn according to a mixture of four Gaussians. The first row of Figure 1 shows the histograms of a sample drawn from this distribution. As similarity function on this data set we choose the Gaussian similarity function $s(x_i, x_j) = \exp(-|x_i - x_j|^2 / 2\sigma^2)$ with $\sigma = 1$. As similarity graph we consider both the fully connected graph and the k-nearest neighbor graph with $k = 10$. In Figure 1 we show the first eigenvalues and eigenvectors of the unnormalized Laplacian L and the normalized Laplacian L_{rw}. That is, in the eigenvalue plot we plot i vs. λ_i (for the moment ignore the dashed line and the different shapes of the eigenvalues in the plots for the unnormalized case; their meaning will be discussed in Section 8.4). In the eigenvector plots of an eigenvector $v = (v_1, \ldots, v_{200})'$ we plot x_i vs. v_i. The first two rows of Figure 1 show the results based on the k-nearest neighbor graph. We can see that the first four eigenvalues are 0, and the corresponding eigenvectors are cluster indicator vectors. The reason is that the clusters form disconnected parts in the k-nearest neighbor graph, in which case the eigenvectors are given as in Propositions 2 and 4. The next two rows show the results for the fully connected graph. As the Gaussian similarity function is always positive, this graph only consists of one connected component. Thus, eigenvalue 0 has multiplicity 1, and the first eigenvector is the constant vector. The following eigenvectors carry the information about the clusters. For example, in the unnormalized case (last row), if we threshold the second eigenvector at 0, then the part below 0 corresponds to clusters 1 and 2, and the part above 0 to clusters 3 and 4. Similarly,

thresholding the third eigenvector separates clusters 1 and 4 from clusters 2 and 3, and thresholding the fourth eigenvector separates clusters 1 and 3 from clusters 2 and 4. Altogether, the first four eigenvectors carry all the information about the four clusters. In all the cases illustrated in this figure, spectral clustering using k-means on the first four eigenvectors easily detects the correct four clusters.

5   Graph cut point of view

The intuition of clustering is to separate points in different groups according to their similarities. For data given in form of a similarity graph, this problem can be restated as follows: we want to find a partition of the graph such that the edges between different groups have a very low weight (which means that points in different clusters are dissimilar from each other) and the edges within a group have high weight (which means that points within the same cluster are similar to each other). In this section we will see how spectral clustering can be derived as an approximation to such graph partitioning problems.

For two disjoint subsets $A, B \subset V$ we define

$$\text{cut}(A, B) = \sum_{i \in A, j \in B} w_{ij}.$$

Given a similarity graph with adjacency matrix W, the simplest and most direct way to construct a partition is to solve the mincut problem. This consists of choosing the partition $A_1, \ldots, A_k$ which minimizes

$$\text{cut}(A_1, \ldots, A_k) := \sum_{i=1}^{k} \text{cut}(A_i, \overline{A}_i)$$

(recall the notation $\overline{A}$ for the complement of a subset $A \subset V$). In particular for $k = 2$, mincut is a relatively easy problem and can be solved efficiently, see Stoer and Wagner (1997) and the discussion therein. However, in practice this often does not lead to satisfactory partitions. The problem is that in many cases, the solution of mincut simply consists in separating one individual vertex from the rest of the graph. Of course this is not what we want to achieve in clustering, as clusters should be reasonably large groups of points. One way to circumvent this problem is to explicitly request that the sets $A_1, ..., A_k$ are "reasonably large". The two most common objective functions which encode this are RatioCut (first introduced by Hagen and Kahng, 1992) and the normalized cut Ncut (first introduced by Shi and Malik, 2000). In RatioCut, the size of a subset A of a graph is measured by its number of vertices $|A|$, while in Ncut the size is measured by the weights of its edges $\text{vol}(A)$. The definitions are:

$$\text{RatioCut}(A_1, \ldots, A_k) = \sum_{i=1}^{k} \frac{\text{cut}(A_i, \overline{A}_i)}{|A_i|}$$

$$\text{Ncut}(A_1, \ldots, A_k) = \sum_{i=1}^{k} \frac{\text{cut}(A_i, \overline{A}_i)}{\text{vol}(A_i)}.$$

Note that both objective functions take a small value if the clusters A_i are not too small. In particular, the minimum of the function $\sum_{i=1}^{k}(1/|A_i|)$ is achieved if all $|A_i|$ coincide, and the minimum of $\sum_{i=1}^{k}(1/\text{vol}(A_i))$ is achieved if all $\text{vol}(A_i)$ coincide. So what both objective functions try to achieve is that the clusters are "balanced", as measured by the number of vertices or edge weights, respectively. Unfortunately, introducing balancing conditions makes the previously simple to solve mincut problem become NP hard, see Wagner and Wagner (1993) for a discussion. Spectral clustering is a way to solve relaxed versions of those problems. We will see that relaxing Ncut leads to normalized spectral clustering, while relaxing RatioCut leads to unnormalized spectral clustering.

5.1   Approximating RatioCut for $k = 2$

Let us start with the case of RatioCut and $k = 2$, because the relaxation is easiest to understand in this setting. Our goal is to solve the optimization problem

$$\min_{A \subset V} \text{RatioCut}(A, \overline{A}). \tag{1}$$

We first will rewrite the problem in a more convenient form. Given a subset $A \subset V$ we define the vector $f = (f_1, \ldots, f_n)' \in \mathbb{R}^n$ with entries

$$f_i = \begin{cases} \sqrt{|\overline{A}|/|A|} & \text{if } v_i \in A \\ -\sqrt{|A|/|\overline{A}|} & \text{if } v_i \in \overline{A}. \end{cases} \tag{2}$$

Now the RatioCut objective function can be conveniently rewritten using the unnormalized graph Laplacian. This is due to the following calculation:

$$
\begin{aligned}
f'Lf &= \sum_{i,j=1}^{n} w_{ij}(f_i - f_j)^2 \\
&= \sum_{i \in A, j \in \overline{A}} w_{ij} \left(\sqrt{\frac{|\overline{A}|}{|A|}} + \sqrt{\frac{|A|}{|\overline{A}|}} \right)^2 + \sum_{i \in \overline{A}, j \in A} w_{ij} \left(-\sqrt{\frac{|\overline{A}|}{|A|}} - \sqrt{\frac{|A|}{|\overline{A}|}} \right)^2 \\
&= 2 \operatorname{cut}(A, \overline{A}) \left(\frac{|\overline{A}|}{|A|} + \frac{|A|}{|\overline{A}|} + 2 \right) \\
&= 2 \operatorname{cut}(A, \overline{A}) \left(\frac{|A| + |\overline{A}|}{|A|} + \frac{|A| + |\overline{A}|}{|\overline{A}|} \right) \\
&= 2|V| \cdot \operatorname{RatioCut}(A, \overline{A}).
\end{aligned}
$$

Additionally, we have

$$
\sum_{i=1}^{n} f_i = \sum_{i \in A} \sqrt{\frac{|\overline{A}|}{|A|}} - \sum_{i \in \overline{A}} \sqrt{\frac{|A|}{|\overline{A}|}} = |A| \sqrt{\frac{|\overline{A}|}{|A|}} - |\overline{A}| \sqrt{\frac{|A|}{|\overline{A}|}} = 0.
$$

In other words, the vector f as defined in Equation (2) is orthogonal to the constant one vector $\mathbb{1}$. Finally, note that f satisfies

$$
\|f\|^2 = \sum_{i=1}^{n} f_i^2 = |A| \frac{|\overline{A}|}{|A|} + |\overline{A}| \frac{|A|}{|\overline{A}|} = |\overline{A}| + |A| = n.
$$

So the problem of minimizing (1) can be equivalently rewritten as

$$\min_{A \subset V} f'Lf \quad \text{subject to } f \perp \mathbb{1}, \ f_i \text{ as defined in Eq. (2)}, \ \|f\| = \sqrt{n}. \tag{3}$$

This is an NP-hard discrete optimization problem as the entries of the solution vector f are only allowed to take two particular values. The obvious relaxation in this setting is to discard the condition on the discrete values for f_i and instead allow $f_i \in \mathbb{R}$. This leads to the relaxed optimization problem

$$\min_{f \in \mathbb{R}^n} f'Lf \quad \text{subject to } f \perp \mathbb{1}, \ \|f\| = \sqrt{n}. \tag{4}$$

By the Rayleigh-Ritz theorem (e.g., see Section 5.5.2. of Lütkepohl, 1997) it can be seen immediately that the solution of this problem is given by the vector f which is the eigenvector corresponding to the second smallest eigenvalue of L (recall that the smallest eigenvalue of L is 0 with eigenvector $\mathbb{1}$). So we can approximate a minimizer of RatioCut by the second eigenvector of L. However, in order to obtain a partition of the graph we need to re-transform the real-valued solution vector f of the relaxed problem into a discrete indicator vector. The simplest way to do this is to use the sign of f as indicator function, that is to choose

$$\begin{cases} v_i \in A & \text{if } f_i \geq 0 \\ v_i \in \overline{A} & \text{if } f_i < 0. \end{cases}$$

However, in particular in the case of $k > 2$ treated below, this heuristic is too simple. What most spectral clustering algorithms do instead is to consider the coordinates f_i as points in $\mathbb{R}$ and cluster them into two groups $C, \overline{C}$ by the k-means clustering algorithm. Then we carry over the resulting clustering to the underlying data points, that is we choose

$$\begin{cases} v_i \in A & \text{if } f_i \in C \\ v_i \in \overline{A} & \text{if } f_i \in \overline{C}. \end{cases}$$

This is exactly the *unnormalized spectral clustering* algorithm for the case of $k = 2$.

5.2   Approximating RatioCut for arbitrary k

The relaxation of the $\mathrm{RatioCut}$ minimization problem in the case of a general value k follows a similar principle as the one above. Given a partition of V into k sets $A_1, \ldots, A_k$, we define k indicator vectors $h_i = (h_{1,i}, \ldots, h_{n,i})'$ by

$$h_{i,j} = \begin{cases} 1/\sqrt{|A_i|} & \text{if } i \in A_j \\ 0 & \text{otherwise.} \end{cases} \tag{5}$$

Then we set the matrix $H \in \mathbb{R}^{n \times k}$ as the matrix containing those k indicator vectors as columns. Observe that the columns in H are orthonormal to each other, that is $H'H = I$. Similar to the calculations in the last section we can see that

$$h_i' L h_i = 2 \frac{\mathrm{cut}(|A_i|, |\overline{A}_i|)}{|A_i|}.$$

Moreover, one can check that

$$h_i' L h_i = (H'LH)_{ii}.$$

Plugging those things together we get

$$\mathrm{RatioCut}(A_1, \ldots, A_k) = \frac{1}{2} \sum_{i=1}^{k} h_i' L h_i = \frac{1}{2} \sum_{i=1}^{k} (H'LH)_{ii} = \frac{1}{2} \mathrm{Tr}(H'LH),$$

where Tr denotes the trace of a matrix. So we can write the problem of minimizing $\mathrm{RatioCut}(A_1, \ldots, A_k)$ as

$$\min_{A_1, \ldots, A_k} \mathrm{Tr}(H'LH) \text{ subject to } H'H = I, \ H \text{ as defined in Eq. (5)}.$$

Similar to above we now relax the problem by allowing the entries of the matrix H to take arbitrary real values. Then the relaxed problem becomes:

$$\min_{H \in \mathbb{R}^{n \times k}} \mathrm{Tr}(H'LH) \text{ subject to } H'H = I.$$

This is the standard form of a trace minimization problem, and again some version of the Rayleigh-Ritz theorem (e.g., see Section 5.2.2.(6) of Lütkepohl, 1997) tells us that the solution is given by choosing H as the matrix which contains the first k eigenvectors of L as columns. We can see that the matrix H is in fact the matrix V used in the unnormalized spectral clustering algorithm as described in Section 4. Again we need to re-convert the real valued solution matrix to a discrete partition. As above, the standard way is to use the k-means algorithms on the rows of V. This then leads to the general unnormalized spectral clustering algorithm as presented in Section 4.

10

5.3 Approximating Ncut

Techniques very similar to the ones used for RatioCut can be used to derive normalized spectral clustering as relaxation of minimizing Ncut. In the case $k = 2$ we define the cluster indicator vector f by

$$f_i = \begin{cases} \sqrt{\frac{\text{vol}(\overline{A})}{\text{vol } A}} & \text{if } i \in A \\ -\sqrt{\frac{\text{vol}(A)}{\text{vol}(\overline{A})}} & \text{if } i \in \overline{A}. \end{cases} \tag{6}$$

Then by similar calculations as above one can check that $(Df)'\mathbb{1} = 0$, $f'Df = \text{vol}(V)$, and $f'Lf = 2\text{vol}(V)\text{Ncut}(A, \overline{A})$. Thus we can rewrite the problem of minimizing Ncut by the equivalent problem

$$\min_A f'Lf \ \text{subject to} \ f \text{ as in (6)}, \ Df \perp \mathbb{1}, \ f'Df = \text{vol}(V). \tag{7}$$

Again we relax the problem by allowing f to be real valued:

$$\min_{f \in \mathbb{R}^n} f'Lf \ \text{subject to} \ Df \perp \mathbb{1}, \ f'Df = \text{vol}(V). \tag{8}$$

Now we substitute $g := D^{1/2}f$. After substitution, the problem is

$$\min_{g \in \mathbb{R}^n} g'D^{-1/2}LD^{-1/2}g \ \text{subject to} \ g \perp D^{1/2}\mathbb{1}, \ \|g\|^2 = \text{vol}(V). \tag{9}$$

Observe that $D^{-1/2}LD^{-1/2} = L_{\text{sym}}$, $D^{1/2}\mathbb{1}$ is the first eigenvector of L_{sym}, and $\text{vol}(V)$ is a constant. Hence, Problem (9) is in the form of the standard Rayleigh-Ritz theorem, and its solution g is given by the second eigenvector of L_{sym}. Re-substituting $f = D^{-1/2}g$ and using Proposition 3 we see that f then is the second eigenvector of L_{rw}, or equivalently the generalized eigenvector of $Lv = \lambda Dv$.

For the case of finding $k > 2$ clusters, we define the indicator vectors $h_i = (h_{1,i}, \ldots, h_{n,i})'$ by

$$h_{i,j} = \begin{cases} 1/\sqrt{\text{vol}(A_i)} & \text{if } i \in A_j \\ 0 & \text{otherwise}. \end{cases} \tag{10}$$

Then we set the matrix H as the matrix containing those k indicator vectors as columns. Observe that $H'H = I$, $h_i'Dh_i = 1$, and $h_i'Lh_i = 2\text{cut}(A_i, \overline{A_i})/\text{vol}(A_i)$. So we can write the problem of minimizing Ncut as

$$\min_{A_1, \ldots, A_k} \text{Tr}(H'LH) \ \text{subject to} \ H'DH = I, \ H \text{ as in (10)}.$$

Relaxing the discreteness condition and substituting $U = D^{1/2}H$ we obtain the relaxed problem

$$\min_{U \in \mathbb{R}^{n \times k}} \text{Tr}(U'D^{-1/2}LD^{-1/2}U) \ \text{subject to} \ U'U = I. \tag{11}$$

Again this is the standard trace minimization problem which is solved by the matrix U which contains the first k eigenvectors of L_{sym} as columns. Re-substituting $H = D^{-1/2}U$ and using Proposition 3 we see that the solution H consists of the first k eigenvectors of the matrix L_{rw}, or the first k generalized eigenvectors of $Lv = \lambda Dv$. This yields the normalized spectral clustering algorithm according to Shi and Malik (2000).

5.4 Comments on the relaxation approach

There are several comments we should make about this derivation of spectral clustering. Most importantly, there is no guarantee whatsoever on the quality of the solution of the relaxed problem compared to the exact solution. That is, if $A_1, \ldots, A_k$ is the exact solution of minimizing RatioCut, and $B_1, \ldots, B_k$ is the solution constructed by unnormalized spectral clustering, then $\text{RatioCut}(B_1, \ldots, B_k) - \text{RatioCut}(A_1, \ldots, A_k)$ can be arbitrary large. An example for the case $k = 2$ can be found in Guattery and Miller (1998). Here the authors consider a very simple class of graphs called "cockroach graphs". Those graphs essentially look like a ladder, with a few rimes removed, see Figure 2. Obviously, the ideal RatioCut just cuts the ladder by a vertical cut such that


Figure 2: The cockroach graph from Guattery and Miller (1998)

$A = \{v_1, \ldots, v_k, v_{2k+1}, \ldots, v_{3k}\}$ and $\overline{A} = \{v_{k+1}, \ldots, v_{2k}, v_{3k+1}, \ldots, v_{4k}\}$. This cut is perfectly balanced with $|A| = |\overline{A}| = 2k$ and $\text{cut}(A, \overline{A}) = 2$. However, by studying the properties of the second eigenvector of the unnormalized graph Laplacian of cockroach graphs the authors prove that unnormalized spectral clustering always cuts horizontally through the ladder, constructing the sets $B = \{v_1, \ldots, v_{2k}\}$ and $\overline{B} = \{v_{2k+1}, \ldots, v_{4k}\}$. This also results in a balanced cut, but now we cut k edges instead of just 2. So $\text{RatioCut}(A, \overline{A}) = 2/k$, while $\text{RatioCut}(B, \overline{B}) = 1$. This means that the RatioCut value of the cut obtained by spectral clustering is $k/2$ times worse than the optimal cut. The same example also works for Ncut. In general it is known that efficient algorithms to approximate balanced graph cuts up to a constant factor do not exist. To the contrary, this approximation problem in NP hard itself (Bui and Jones, 1992).

Of course, the relaxation we discussed above is not unique. For example, a completely different relaxation which leads to a semi-definite program is derived in Bie and Cristianini (2006), and there might be many other useful relaxations. The reason why the spectral relaxation is so appealing is not that it leads to particularly good solutions, but is the fact that it results in a very simple to solve standard linear algebra problem.

Finally, there is nothing principled about using the k-means algorithm to construct discrete partitions from the real valued representation vectors y_i. Any other algorithm which can solve this problem could be used instead, and several other techniques are regularly used. For example, we can try to construct the partition by separating the points y_i by hyperplanes in $\mathbb{R}^k$, see Lang (2006) for a detailed discussion of this technique. However, one can argue that at least the Euclidean distance between the points y_i is a meaningful quantity to look at. In the next section we will see that the Euclidean distance between the points y_i is related to the "commute distance" on the graph. In Nadler, Lafon, Coifman, and Kevrekidis (2006), the authors show that the Euclidean distances between the y_i are also related to a more general "diffusion distance". The spectral embedding also turns out to be a solution of several optimization problems on the graph, for further examples see Bolla (1991) or Belkin and Niyogi (2003). In all those optimization problems, the Euclidean distance in the embedding space turns out to be a meaningful quantity. Bach and Jordan (2004) use a more advanced post-processing of the eigenvectors. They study the subspace spanned by the first k eigenvectors, and try to approximate this subspace as good as possible using piecewise constant vectors. This also leads to minimizing certain Euclidean distances in the space $\mathbb{R}^k$, which can be done by some weighted k-means algorithm.

6 Random walks point of view

Another line of argument to explain spectral clustering is based on random walks on the similarity graph. A random walk on a graph is a stochastic process which randomly jumps from vertex to vertex. We will see below that spectral clustering can be interpreted as trying to find a partition of the graph such that the random walk stays

long within the same cluster and seldom jumps between clusters. Intuitively this makes sense, in particular together with the graph cut explanation of the last section. A partition with a low cut will also have the property that the random walk does not have many opportunities to jump between clusters. Formally, the transition probability of jumping in one step from vertex i to vertex j is proportional to the edge weight w_{ij} and is given by $p_{ij} := w_{ij}/d_i$. The transition matrix $P = (p_{ij})_{i,j=1,...,n}$ of the random walk is thus defined by

$$P = D^{-1}W.$$

If the graph is connected and non-bipartite, then the random walk always possesses a unique stationary distribution $\pi = (\pi_1, \ldots, \pi_n)'$, which is given by $\pi_i = d_i/\operatorname{vol}(G)$. For background reading on random walks in general we refer to Norris (1997) and Brémaud (1999), and for random walks on graphs we recommend Aldous and Fill (in preparation) and Lovász (1993).

Obviously there is a tight relationship between L_{rw} and P, as $L_{\mathrm{rw}} = I - P$. As a consequence, λ is an eigenvalue of L_{rw} with eigenvector v if and only if $1 - \lambda$ is an eigenvector of P with eigenvector v. It is well known that many properties of a graph can be expressed in terms of the matrix P, see Lovász (1993) for an overview. So from a random walk point of view it does not come as a surprise that the largest eigenvectors of P and the smallest eigenvectors of L_{rw} can be used to describe cluster properties of the graph.

Random walks and Ncut

A formal equivalence between Ncut and transition probabilities of the random walk has been observed in Meila and Shi (2001).

Proposition 5 (Ncut **via transition probabilities**) *Let G be connected and not bi-partite. Assume that we run the random walk $(X_t)_{t \geq 0}$ starting with X_0 in the stationary distribution π. For disjoint subsets $A, B \subset V$, denote by $P(B|A) := P(X_1 \in B | X_0 \in A)$. Then:*

$$\mathrm{Ncut}(A, \overline{A}) = P(\overline{A}|A) + P(A|\overline{A}).$$

Proof. First of all observe that

$$P(X_0 \in A, X_1 \in B) = \sum_{i \in A, j \in B} P(X_0 = i, X_1 = j) = \sum_{i \in A, j \in B} \pi_i p_{ij}$$

$$= \sum_{i \in A, j \in B} \frac{d_i}{\operatorname{vol}(G)} \frac{w_{ij}}{d_i} = \frac{1}{\operatorname{vol}(G)} \sum_{i \in A, j \in B} w_{ij}.$$

Using this we obtain

$$P(X_1 \in B | X_0 \in A) = \frac{P(X_0 \in A, X_1 \in B)}{P(X_0 \in A)}$$

$$= \left(\frac{1}{\operatorname{vol}(G)} \sum_{i \in A, j \in B} w_{ij} \right) \left(\frac{\operatorname{vol}(A)}{\operatorname{vol}(G)} \right)^{-1} = \frac{\sum_{i \in A, j \in B} w_{ij}}{\operatorname{vol}(A)}.$$

Now the proposition follows directly with the definition of Ncut. ☺

This proposition leads to a nice interpretation of Ncut, and hence of normalized spectral clustering. It tells us that when minimizing Ncut, we actually look for a cut through the graph such that a random walk seldom transitions from A to $\overline{A}$ or vice versa.

The commute distance

A second tight connection between random walks and graph Laplacians can be made via the commute distance on the graph. The commute distance (also called resistance distance) $c(i, j)$ between two vertices i and j is the expected time it takes the random walk to travel from vertex i to vertex j and back (Lovász, 1993; Aldous and Fill, in preparation). The commute distance has several nice properties which make it particularly appealing for

machine learning. As opposed to the shortest path distance on a graph, the commute distance between two vertices decreases if there are many different short ways to get from vertex i to vertex j. So instead of just looking for the one shortest path, the commute distance looks at the set of short paths. Points which are connected by a short path and lie in the same cluster of the graph are much closer to each other than points which are connected by a short part but lie in different clusters.

Remarkably, the commute distance on a graph can be computed with the help of the generalized inverse (also called pseudo-inverse or Moore-Penrose inverse) $L^\dagger$ of the graph Laplacian L. To define the generalized inverse of L, recall that by Proposition 1, the matrix L can be decomposed as $L = V \Lambda V'$ where V is the matrix containing the eigenvectors as columns and Λ the diagonal matrix with the eigenvalues $\lambda_1, \ldots, \lambda_n$ on the diagonal. As at least one of the eigenvalues is 0, the matrix L is not invertible. Instead, we define its generalized inverse as $L^\dagger := V \Lambda^\dagger V'$ where the matrix $\Lambda^\dagger$ is the diagonal matrix with diagonal entries $1/\lambda_i$ if $\lambda_i \neq 0$ and 0 if $\lambda_i = 0$. The entries of $L^\dagger$ can be computed as $l_{ij}^\dagger = \sum_{k=2}^n \frac{1}{\lambda_k} v_{ik} v_{jk}$. Moreover, as all its eigenvalues are non-negative, $L^\dagger$ is positive semi-definite. For further properties of $L^\dagger$ see Gutman and Xiao (2004).

Proposition 6 (Commute distance) *Let $G = (V, E)$ a connected, undirected graph. Denote by c_{ij} the commute distance between vertex i and vertex j, and by $L^\dagger = (l_{ij}^\dagger)_{i,j=1,\ldots,n}$ the generalized inverse of L. Then we have:*

$$c_{ij} = \mathrm{vol}(G)(l_{ii}^\dagger - 2l_{ij}^\dagger + l_{jj}^\dagger) = \mathrm{vol}(G)(e_i - e_j)'L^\dagger(e_i - e_j).$$

This result has been published by Klein and Randic (1993), where it has been proved by methods of electrical network theory. For a proof using first step analysis for random walks see Fouss, Pirotte, Renders, and Saerens (2006). As both proofs are technical we omit them.

There also exist other ways to express the commute distance with the help of graph Laplacians. For example a method in terms of eigenvectors of the normalized Laplacian L_{sym} can be found as Corollary 3.2 in Lovász (1993), and a method computing the commute distance with the help of determinants of certain sub-matrices of L can be found in Bapat, Gutman, and Xiao (2003).

Proposition 6 has an important consequence. It shows that $\sqrt{c_{ij}}$ can be considered as a Euclidean distance function on the vertices of the graph. This means that we can construct an embedding which maps the vertices v_i of the graph on points $z_i \in \mathbb{R}^n$ such that the Euclidean distances between the points z_i coincide with the commute distances on the graph. This works as follows. As the matrix $L^\dagger$ is positive semi-definite, it induces an inner product on $\mathbb{R}^n$ (or to be more formal, it induces an inner product on the subspace of $\mathbb{R}^n$ which is perpendicular to the vector $\mathbb{1}$). Now choose z_i as the point in $\mathbb{R}^n$ corresponding to the i-th row of the matrix $(\Lambda^\dagger)^{1/2}V$. Then, by Proposition 6 and by the construction of $L^\dagger$ we have that $\langle z_i, z_j \rangle = e_i'L^\dagger e_j$ and $||z_i - z_j||^2 = c_{ij}$.

The embedding used in unnormalized spectral clustering is related to the commute time embedding, but not identical. In spectral clustering, we map the vertices of the graph on the rows y_i of the matrix V, while the commute time embedding maps the vertices on the rows z_i of the matrix $(\Lambda^\dagger)^{1/2}V$. That is, compared to the entries of y_i, the entries of z_i are additionally scaled by the inverse eigenvalues of L. Moreover, in spectral clustering we only take the first k columns of the matrix, while the commute time embedding takes all columns of the matrix. Several authors now try to justify why y_i and z_i are not so different after all, and then state a bit hand-waiving that the fact that spectral clustering constructs clusters based on the Euclidean distances between the y_i can be interpreted as building clusters of the vertices in the graph based on the commute distance. However, while this intuition might be helpful, we have not seen a precise derivation of such a claim.

7 Perturbation theory point of view

In Section 3 we have already seen that if the graph consists of k disconnected components, then the multiplicity of the eigenvalue 0 of both L and L_{rw} is k, and the eigenspace is spanned by the indicator vectors of the connected components. Briefly stated, the perturbation argument now says that if we do not have a completely ideal situation where the between-cluster similarity is exactly 0, but if we have a situation where the between-cluster similarities are very small, then the eigenvectors of the first k eigenvalues should be very close to the ones in the ideal case. Thus we should still be able to recover the clustering from those eigenvectors.

Perturbation theory studies the question how eigenvalues and eigenvectors of a matrix A change if we add a small perturbation H, that is we consider the perturbed matrix $\tilde{A} := A + H$. Most perturbation theorems state that a certain distance between eigenvalues or eigenvectors of A and $\tilde{A}$ is bounded by a constant times a norm of H. The constant usually depends on which eigenvalue we are looking at, and how far this eigenvalue is separated from the rest of the spectrum (for a formal statement see below). The justification of spectral clustering is then the following: Let us first consider the "ideal case" where the between-cluster similarity is 0. Here, the first k eigenvectors of L or L_{rw} are the indicator vectors of the clusters. In this case, the points $y_i \in \mathbb{R}^k$ constructed in the spectral clustering algorithms have the form $(0, \ldots, 0, 1, 0, \ldots 0)'$ where the position of the 1 indicates the connected component this point belongs to. In particular, all y_i belonging to the same connected component coincide. The k-means algorithm will trivially find the correct partition by placing a center point on each of the points $(0, \ldots, 0, 1, 0, \ldots 0)' \in \mathbb{R}^k$. In a "nearly ideal case" where we still have distinct clusters, but the between-cluster similarity is not exactly 0, we consider the Laplacian matrices to be perturbed versions of the ones of the ideal case. Perturbation theory then tells us that the eigenvectors will be very close to the ideal indicator vectors. The points y_i might not completely coincide with $(0, \ldots, 0, 1, 0, \ldots 0)'$, but do so up to some small error term. Hence, if the perturbations are not too large, then k-means algorithm will still separate the groups from each other.

Formally, those arguments are based on the Davis-Kahan theorem from matrix perturbation theory, which bounds the difference between eigenspaces of symmetric matrices under perturbations. We state those results for completeness, but for background reading we refer to Section V of Stewart and Sun (1990) and Section VII.3 of Bhatia (1997). In perturbation theory, distances between subspaces are usually measured using "canonical angles" (also called "principal angles"). To define principal angles, let $\mathcal{V}_1$ and $\mathcal{V}_2$ be two p-dimensional subspaces of $\mathbb{R}^d$, and V_1 and V_2 two matrices such that their columns form orthonormal systems for $\mathcal{V}_1$ and $\mathcal{V}_2$, respectively. Then the cosines $\cos \Theta_i$ of the principal angles Θ_i are the singular values of $V_1' V_2$. For $p = 1$, the so defined canonical angles coincide with the normal definition of an angle. Canonical angles can also be defined if $\mathcal{V}_1$ and $\mathcal{V}_2$ do not have the same dimension, see Section V of Stewart and Sun (1990), Section VII.3 of Bhatia (1997), or Section 12.4.3 of Golub and Van Loan (1996). The matrix $\sin \Theta(\mathcal{V}_1, \mathcal{V}_2)$ will denote the diagonal matrix with the canonical angles on the diagonal.

Theorem 7 (Davis-Kahan) *Let $A, H \in \mathbb{R}^{n \times n}$ be symmetric matrices, and let $\| \cdot \|$ be the Frobenius norm or the two-norm for matrices. Consider $\tilde{A} := A + H$ as a perturbed version of A. Let $S_1 \subset \mathbb{R}$ be an interval. Denote by $\sigma_{S_1}(A)$ the set of eigenvalues of A which are contained in S_1, and by V_1 the eigenspace corresponding to all those eigenvalues (more formally, V_1 is the image of the spectral projection induced by $\sigma_{S_1}(A)$). Denote by $\sigma_{S_1}(\tilde{A})$ and $\tilde{V}_1$ the analogous quantities for $\tilde{A}$. Define the distance between S_1 and the spectrum of A outside of S_1 as*

$$\delta = \min\{|\lambda - s|; \ \lambda \ eigenvalue \ of \ A, \ \lambda \notin S_1, \ s \in S_1\}.$$

Then the distance $d(V_1, \tilde{V}_1) := \| \sin \Theta(V_1, \tilde{V}_1) \|$ between the two subspaces V_1 and $\tilde{V}_1$ is bounded by

$$d(V_1, \tilde{V}_1) \leq \frac{\|H\|}{\delta}.$$

Let us try to decrypt this theorem, for simplicity in the case of the unnormalized Laplacian (for the normalized Laplacian it works analogously). The matrix A will correspond to the graph Laplacian L in the ideal case, that is we assume that the graph has k connected components. The matrix $\tilde{A}$ corresponds to a perturbed case, where due to noise the k components in the graph are no longer completely disconnected, but they are only connected by few edges with low weight. We denote the corresponding graph Laplacian of this case by $\tilde{L}$. For spectral clustering, we need to consider the first k eigenvalues and eigenvectors of $\tilde{L}$. Denote the eigenvalues of L by $\lambda_1, \ldots \lambda_n$ and the ones of the perturbed Laplacian $\tilde{L}$ by $\tilde{\lambda}_1, \ldots, \tilde{\lambda}_n$. Choosing the interval S_1 is now the crucial point. We want to choose it such that both the first k eigenvalues of $\tilde{L}$ and the first k eigenvalues of L are contained in S_1. This is easier the smaller the perturbation $H = L - \tilde{L}$ and the larger the eigengap $|\lambda_k - \lambda_{k+1}|$ is. If we manage to find such a set, then the Davis-Kahan theorem tells us that the eigenspaces corresponding to the first k eigenvalues in ideal case L and the first k eigenvalues in the perturbed case $\tilde{L}$ are very close to each other, that is their distance is bounded by $\|H\|/\delta$. Then, as the eigenvectors in the ideal case are piecewise constant on the connected components, this will approximately also be true in the perturbed case. How good "approximately" is depends on the norm of the perturbation $\|H\|$ and the distance δ between S_1 and the $(k + 1)$st eigenvector of

L. In case the set S_1 has been chosen as the interval $[0, \lambda_k]$, then δ coincides with the spectral gap $|\lambda_{k+1} - \lambda_k|$. We can see from the theorem that the larger this eigengap is, the closer the eigenvectors of the ideal case and the perturbed case are, and hence the better spectral clustering works. Later we will see that the size of the eigengap as a quality criterion for spectral clustering also emerges from the other explanations of spectral clustering.

If the perturbation H is too large or the eigengap is too small, we might not find a set S_1 such that both the first k eigenvalues of L and $\tilde{L}$ are contained. In this case, we need to make a compromise by choosing the set to contain the first k eigenvalues of L, but maybe a bit more or fewer eigenvalues of $\tilde{L}$. The statement of the theorem then becomes weaker.

7.1   Comments about the perturbation approach

A bit of caution is needed when using perturbation theory arguments to justify clustering algorithms based on eigenvectors of matrices. In general, *any* symmetric block diagonal symmetric matrix has the property that there exists a basis of eigenvectors which are zero outside the individual blocks and real-valued within the blocks. For example, based on this argument several authors use the eigenvectors of the similarity matrix S or adjacency matrix W to discover clusters. However, being block diagonal in the ideal case of completely separated clusters can be considered as a necessary condition for a successful use of eigenvectors, but not a sufficient one. At least two more properties should be satisfied:

First, we need to make sure that the *order* of the eigenvalues and eigenvectors is meaningful. In case of the Laplacians this is always true, as we know that any connected component possesses exactly one eigenvector which has eigenvalue 0. Hence, if the graph has k connected components and we take the first k eigenvectors of the Laplacian, then we know that we have exactly one eigenvector per component. However, this might not be the case for other matrices such as S or W. For example, it could be the case that the two largest eigenvalues of a block diagonal similarity matrix S come from the same block. In such a situation, if we take the first k eigenvectors of S, some blocks will be represented several times, while there are other blocks which we will miss completely (unless we take certain precautions). This is the reason why using the eigenvectors of S or W for clustering should be discouraged.

The second property is that in the ideal case, the entries of the eigenvectors on the components should be "safely bounded away" from 0. Assume that an eigenvector on the first connected component has an entry $v_{1,i} = \varepsilon$ at position i. In the ideal case, the fact that this entry is non-zero indicates that the corresponding point i belongs to the first cluster. The other way round, if a point j does not belong to cluster 1, then in the ideal case it should be the case that $v_{1,j} = 0$. Now consider the same situation, but with perturbed data. The perturbed eigenvector $\tilde{v}$ will usually not have any non-zero component any more; but if the noise is not too large, then perturbation theory tells us that the entries $\tilde{v}_{1,i}$ and $\tilde{v}_{1,j}$ are still "close" to their original values $v_{1,i}$ and $v_{1,j}$. So both entries $\tilde{v}_{1,i}$ and $\tilde{v}_{1,j}$ will take some small values, say ε_1 and ε_2. In practice, it is now very unclear how we should interpret this situation. Either we believe that small entries in $\tilde{v}$ indicate that the points do not belong to the first cluster (which then misclassifies the first data point i), or we think that the entries already indicate class membership and classify both points to the first cluster (which misclassifies point j).

For both matrices L and L_{rw}, the eigenvectors in the ideal situation are indicator vectors, so the problems described above cannot occur. However, this is not true for the matrix L_{sym}, which is used in the normalized spectral clustering algorithm of Ng et al. (2002). Even in the ideal case, the eigenvectors of this matrix are given as $D^{1/2}\mathbb{1}_{A_i}$. If the degrees of the vertices differ a lot, and in particular if there are vertices which have a very low degree, the corresponding entries in the eigenvectors are very small. To counteract the problem described above, the row-normalization step in the algorithm of Ng et al. (2002) comes into play. In the ideal case, the matrix V in the algorithm has exactly one non-zero entry per row. After row-normalization, the matrix U in the algorithm of Ng et al. (2002) thus consists of the cluster indicator vectors. Note however, that in practice different things might happen. Assume that we have $\tilde{v}_{1,i} = \varepsilon_1$ and $\tilde{v}_{2,i} = \varepsilon_2$. If we now normalize the i-th row of V, both ε_1 and ε_2 will be multiplied by the factor of $1/\sqrt{\varepsilon_1^2 + \varepsilon_2^2}$ and become rather large. We now run into a similar problem as described above: both points are likely to be classified into the same cluster, even though they belong to different clusters. This argument shows that spectral clustering using the matrix L_{sym} can be problematic if the eigenvectors contain particularly small entries. However, note that such small entries in the eigenvectors only occur if some of the vertices have a particularly low degrees (as the eigenvectors of L_{sym} are given by $D^{1/2}\mathbb{1}_{A_i}$). One could argue

that in such a case, the data point should be considered an outlier anyway, and then it does not really matter in which cluster the point will end up.

To summarize, the conclusion is that both unnormalized spectral clustering and normalized spectral clustering with L_{rw} are well justified by the perturbation theory approach. Normalized spectral clustering with L_{sym} can also be justified by perturbation theory, but it should be treated with more care if the graph contains vertices with very low degrees.

8   Practical details

In this section we will briefly discuss some of the issues which come up when actually implementing spectral clustering. There are several choices to be made and parameters to be set. However, the short discussion in this section is mainly meant to raise awareness about the general problems which can occur. We will look at toy examples only. For thorough studies on the behavior of spectral clustering for various real world tasks we refer to the literature.

8.1   Constructing the similarity graph

Choosing the similarity graph and its parameters for spectral clustering is not a trivial task. This already starts with the choice of the similarity function s_{ij} itself. In general one should try to ensure that the local neighborhoods induced by this similarity function are "meaningful", but in particular in a clustering setting this is very difficult to assess. Ultimately, the choice of the similarity function depends on the domain the data comes from, and no general rules can be given. The second choice concerns the construction of the similarity graph, that is which type of graph we choose and how we set the parameter which governs its connectedness (e.g., the parameter ε of the ε-neighborhood graph or the parameter k of the k-nearest neighbor graph).
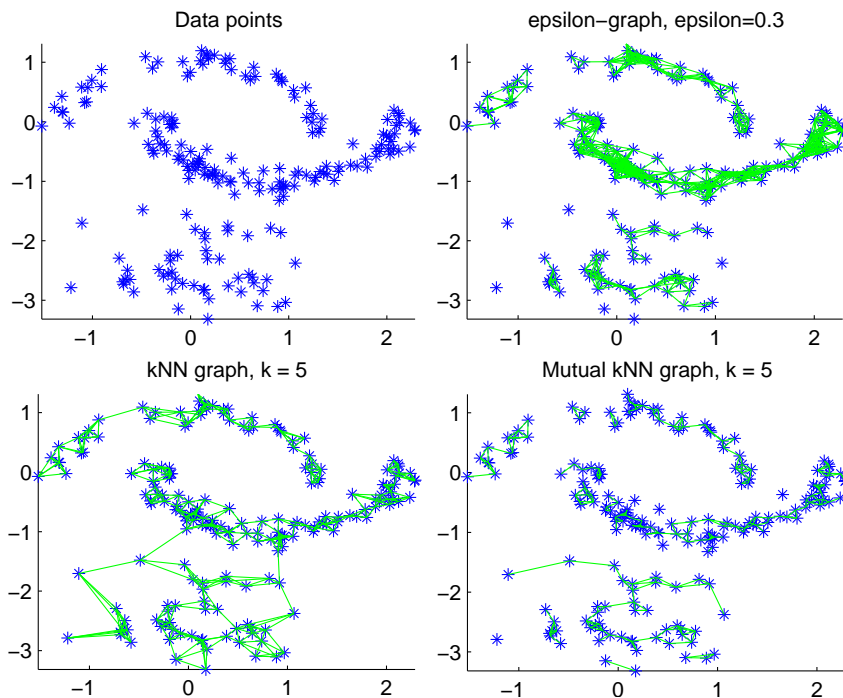

Figure 3: Different similarity graphs, see text for details.

To illustrate the behavior of the different graphs we use the toy example presented in Figure 3. As underlying distribution we choose a distribution on $\mathbb{R}^2$ with three clusters: two "moons" and a Gaussian. The density of the bottom moon is chosen to be larger than the one of the top moon. The upper left panel in Figure 3 shows a sample

drawn from this distribution. The next three panels show the different similarity graphs on this sample. In the ε-neighborhood graph, we can see that it is difficult to choose a useful parameter ε. With $\varepsilon = 0.3$ as in the figure, the points on the middle moon are already very tightly connected, while the points in the Gaussian are barely connected. This problem always occurs if we have data "on different scales", that is the distances between data points are different in different regions of the space.

The k-nearest neighbor graph, on the other hand, can connect points "on different scales". We can see that points in the low-density Gaussian are connected with points in the high-density moon. This is a general property of k-nearest neighbor graphs which can be very useful. We can also see that the k-nearest neighbor graph can fall into several disconnected components if there are high density regions which are reasonably far away from each other. This is the case for the two moons in this example.

The mutual k-nearest neighbor graph has the property that it tends to connect points within regions of constant density, but does not connect regions of different densities with each other. So the mutual k-nearest neighbor can be considered as being "in between" the ε-neighborhood graph and the k-nearest neighbor graph. It is able to act on different scales, but does not mix those scales with each other. Hence, the mutual k-nearest neighbor graph seems particularly well-suited if we want to detect clusters of different densities.

The fully connected graph is very often used in connection with the Gaussian similarity function $s(x_i, x_j) = \exp(-\|x_i - x_j\|^2 / 2\sigma^2)$. Here the parameter σ plays a similar role as the parameter ε in the ε-neighborhoods. Points in local neighborhoods are connected with a relatively high weights, while edges between far away points have a positive, but negligible weights. However, the resulting similarity matrix is not a sparse matrix.

In general, spectral clustering can be quite sensitive to changes in the similarity graph and to the choice of its parameters. Unfortunately, to our knowledge there has been no systematic study which investigates the effects of the similarity graph on clustering and comes up with simple rules of thumb.

8.2 Computing the eigenvectors

To implement spectral clustering in practice one has to compute the first k eigenvectors of a potentially large Laplace matrix. Luckily, if we use the k-nearest neighbor graph or the ε-neighborhood graph, then all Laplace matrices are sparse. Efficient methods exist to compute the first eigenvectors of sparse matrices, the most popular ones being the power method or Krylov subspace methods such as the Lanczos method (Golub and Van Loan, 1996). The speed of convergence of those algorithms depends on the size of the eigengap (also called spectral gap) $\gamma_k = |\lambda_k - \lambda_{k+1}|$. The larger this eigengap is, the faster the algorithms computing the first k eigenvectors converge.

Note that a general problem occurs if one of the eigenvalues under consideration has multiplicity larger than one. For example, in the ideal situation of k disconnected clusters, the eigenvalue 0 has multiplicity k. As we have seen, in this case the eigenspace is spanned by the k cluster indicator vectors. But unfortunately, the vectors computed by the numerical eigensolvers do not necessarily converge to those particular vectors. Instead they just converge to some orthonormal basis of the eigenspace, and it usually depends on implementation details to which basis exactly the algorithm converges. But this is not so bad after all. Note that all vectors in the space spanned by the cluster indicator vectors $\mathbb{1}_{A_i}$ have the form $v = \sum_{i=1}^{k} a_i \mathbb{1}_{A_i}$ for some coefficients a_i, that is, they are piecewise constant on the clusters. So the vectors returned by the eigensolvers still encode the information about the clusters, which can then be used by the k-means algorithm to reconstruct the clusters.

8.3 The number of clusters

Choosing the number k of clusters is a general problem for all clustering algorithms, and a variety of more or less successful methods have been devised for this problem, for example the gap statistic in Tibshirani, Walther, and Hastie (2001) (note that this gap has nothing to do with the eigengap), or stability approaches, see Ben-Hur, Elisseeff, and Guyon (2002), Lange, Roth, Braun, and Buhmann (2004), but also Ben-David, von Luxburg, and Pal (2006). One tool which is particularly designed for spectral clustering is the eigengap heuristic, which can be used for all three graph Laplacians. Here the goal is to choose the number k such that all eigenvalues $\lambda_1, \ldots, \lambda_k$ are very small, but λ_{k+1} is relatively large. There are several justifications for this procedure. The first one is based on perturbation theory, where we observe that in the ideal case of k completely disconnected clusters, the

eigenvalue 0 has multiplicity k, and then there is a gap to the $(k + 1)$th eigenvalue $\lambda_{k+1} > 0$. Other explanations can be given by spectral graph theory. Here, many geometric invariants of the graph can be expressed or bounded with the help of the first eigenvalues of the graph Laplacian. In particular, the sizes of cuts are closely related to the size of the first eigenvalues. For more details on this topic we refer to (Mohar, 1997) and Chung (1997).

We would like to illustrate the eigengap heuristic on our toy example introduced in Section 4. For this purpose we consider similar data sets as in Section 4, but to vary the difficulty of clustering we consider the Gaussians with increasing variance. The first row of Figure 4 shows the histograms of the three samples. We construct the 10-nearest neighbor graph as described in Section 4, and plot the eigenvalues and eigenvectors of the normalized Laplacian L_{rw} on the different samples (the results for the unnormalized Laplacian are similar). The first data set consists of four well separated clusters, and we can see that the first 4 eigenvalues are approximately 0. Then there is a gap between the 4th and 5th eigenvalue, that $|\lambda_5 - \lambda_4|$ is relatively large. According to the eigengap heuristic, this gap indicates that the data set contains 4 clusters. The same behavior can also be observed for the results of the fully connected graph (already plotted in Figure 1). So we can see that the heuristic works well if the clusters in the data are very well pronounced. However, the more noisy or overlapping the clusters are, the less effective is this heuristic. We can see that for the second data set where the clusters are more "blurry", there is still a gap between the 4th and 5th eigenvalue, but it is not so clear to detect as in the case before. Finally, in the last data set, there is no well-defined gap, as the differences between all eigenvalues are approximately the same. But on the other hand, the clusters in this data set overlap so much that one might as well state that there are no clear clusters in the data anyway. This illustrates that, as most methods for choosing the number of clusters, the eigengap heuristic usually works well if the data contains very well pronounced clusters, but in ambiguous cases it also returns ambiguous results.
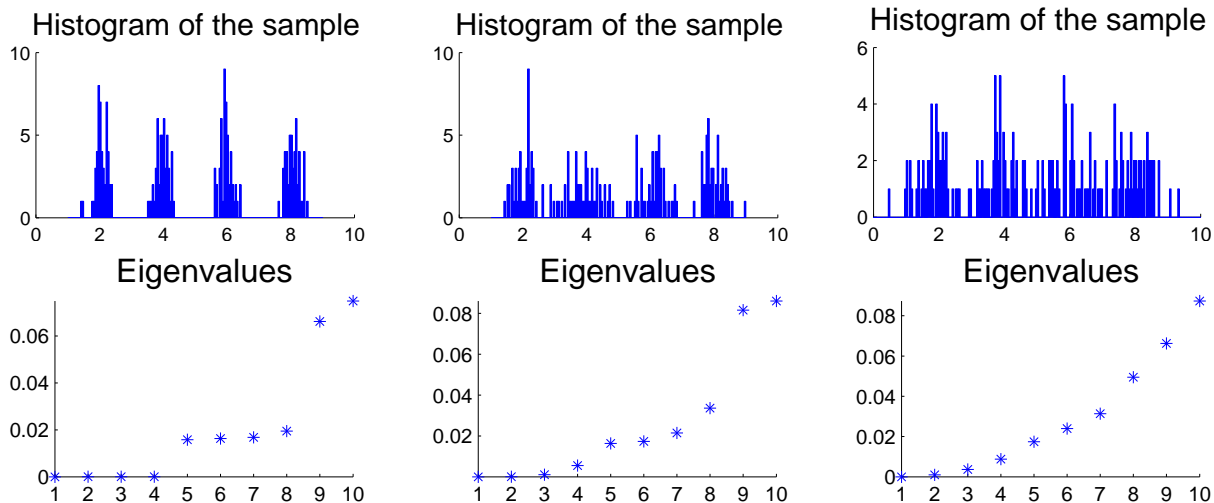

Figure 4: Three data sets, and the smallest 10 eigenvalues of L_{rw}.

8.4   Which graph Laplacian should be used?

A fundamental question related to spectral clustering is the question which of the three graph Laplacians should be used to compute the eigenvectors. Before deciding this question, one should always look at the degree distribution of the similarity graph. If the graph is very regular and most vertices have approximately the same degree, then all the Laplacians are very similar to each other, and will work equally well for clustering. However, if the degrees in the graph are very broadly distributed, then the Laplacians differ considerably. In our opinion, there are several arguments which advocate for using normalized rather than unnormalized spectral clustering, and in the normalized case to use the eigenvectors of L_{rw} rather than those of L_{sym}.

Clustering objectives satisfied by the different algorithms

The first argument in favor of normalized spectral clustering comes from the graph partitioning point of view. For simplicity let us discuss the case $k = 2$. In general, clustering has two different objectives:

1. We want to find a partition such that points in different clusters are dissimilar to each other, that is we want to minimize the between-cluster similarity. In the graph setting, this means to minimize $\sum_{i \in A, j \in \overline{A}} w_{ij}$.

2. We want to find a partition such that points in the same cluster are similar to each other, that is we want to maximize the within-cluster similarity. This means that $\sum_{i,j \in A} w_{ij}$ and $\sum_{i,j \in \overline{A}} w_{ij}$ should be maximized.

Both $\mathrm{RatioCut}$ and Ncut directly implement the first point by explicitly incorporating $\mathrm{cut}(A, \overline{A})$ in the objective function. However, concerning the second point, both algorithms behave differently. Note that

$$\sum_{i,j \in A} w_{ij} = \sum_{i \in A, j \in A \cup \overline{A}} w_{ij} - \sum_{i \in A, j \in \overline{A}} w_{ij} = \mathrm{vol}(A) - \mathrm{cut}(A, \overline{A}).$$

Hence, the within-cluster similarity is maximized if $\mathrm{cut}(A, \overline{A})$ is small *and* if $\mathrm{vol}(A)$ is large. In the case of Ncut, this is also part of the objective function as we want to maximize both $\mathrm{vol}(A)$ and $\mathrm{vol}(\overline{A})$. Thus, Ncut implements the second objective. This can be seen even more explicitly by considering yet another graph cut objective function, namely the $\mathrm{MinMaxCut}$ criterion introduced by Ding, He, Zha, Gu, and Simon (2001):

$$\mathrm{MinMaxCut}(A, B) := \sum_{i=1}^{k} \frac{\mathrm{cut}(A_i, \overline{A}_i)}{\sum_{i,j \in A} w_{ij}}.$$

Here the denominator directly contains the within-cluster similarity, rather than the sum of the within-cluster similarity and the cut, as it is the case for Ncut. But as a good Ncut solution will have a small value of $\mathrm{cut}(A, \overline{A})$ anyway, Ncut and $\mathrm{MinMaxCut}$ are often minimized by similar cuts. Moreover, relaxing $\mathrm{MinMaxCut}$ leads to exactly the same optimization problem as relaxing Ncut, namely to normalized spectral clustering with the eigenvectors of L_{rw}.

Now consider the case of $\mathrm{RatioCut}$. Here the objective is to maximize $|A|$ and $|\overline{A}|$ instead of $\mathrm{vol}(A)$ and $\mathrm{vol}(\overline{A})$. But $|A|$ and $|\overline{A}|$ are not necessarily related to the within-cluster similarity, as the within-cluster similarity depends on the edges and not on the number of vertices in A. Just imagine a set A which has very many vertices, all of which only have very low weighted edges to each other. So minimizing $\mathrm{RatioCut}$ does not attempt to maximize the within-cluster similarity, and the same is then true for its relaxation by unnormalized spectral clustering.

So this is our first important point to keep in mind: Normalized spectral clustering implements both clustering objectives mentioned above, while unnormalized spectral clustering only implements the first objective.

Consistency issues

A completely different argument for the superiority of normalized spectral clustering comes from a statistical analysis of both algorithms. It can be proved (von Luxburg, Bousquet, and Belkin, 2004, 2005; von Luxburg, Belkin, and Bousquet, 2004) that under very mild conditions, both normalized spectral clustering algorithms are statistically consistent. This means that if we assume that the data has been sampled randomly according to some probability distribution from some underlying space, and if we let the sample size increase to infinity, then the results of normalized spectral clustering converge, and the limit partition is usually a sensible partition of the underlying space. Those results do not necessarily hold for unnormalized spectral clustering. It can be proved that unnormalized spectral clustering can fail to converge, or that it can converge to trivial solutions which construct clusters consisting of one single point of the data space.

There is a simple necessary condition which has to be satisfied to avoid such trivial solutions: the eigenvalues of L corresponding to the eigenvectors used in unnormalized spectral clustering must be significantly below the minimal degree in the graph. This means that if we use the first k eigenvectors, then $\lambda_i \ll \min_{j=1,\ldots,n} d_j$

Figure 5: Consistency of unnormalized spectral clustering. Plotted are eigenvalues and eigenvectors of L, for parameter $\sigma = 2$ (first row) and $\sigma = 5$ (second row). The dashed line indicates $\min d_j$, the eigenvalues below $\min d_j$ are plotted as red diamonds, the eigenvalues above $\min d_j$ are plotted as blue stars.

should hold for all $i = 1, \ldots, k$. The reason is that eigenvectors corresponding to eigenvalues with $\lambda \geq \min d_j$ approximate Dirac functions, that is they are approximately 0 in all but one coordinate. If those eigenvectors are used for clustering, then they separate the one vertex where the eigenvector is non-zero from all other vertices, and we clearly do not want to construct such a partition. As an illustration of this phenomenon, consider again our toy data set from Section 4. We consider the first eigenvalues and eigenvectors of the unnormalized graph Laplacian based on the fully connected graph, for different choices of the parameter σ of the Gaussian similarity function (see last row of Figure 1 and Figure 5). The eigenvalues above $\min d_j$ are plotted as blue stars, the eigenvalues below $\min d_j$ are plotted as red diamonds. The dashed line indicates $\min d_j$. In general, we can see that the eigenvectors corresponding to eigenvalues which are much below the dashed lines are "useful" eigenvectors. In case $\sigma = 1$ (plotted already in the last row of Figure 1), Eigenvalues 2, 3 and 4 are significantly below $\min d_j$, and the corresponding Eigenvectors 2, 3, and 4 are meaningful (as already discussed in Section 4. If we increase the parameter σ, we can observe that the eigenvalues tend to move towards $\min d_j$. In case $\sigma = 2$, only the first three eigenvalues are below $\min d_j$ (first row in Figure 5), and in case $\sigma = 5$ only the first two eigenvalues are below $\min d_j$ (second row in Figure 5). We can see that as soon as an eigenvalue gets close to or above $\min d_j$, its corresponding eigenvector approximates a Dirac function. Of course, those eigenvectors are unsuitable for constructing a clustering. We would like to stress that those problems only concern the eigenvectors of the matrix L, and they do not occur for L_{rw} or L_{sym}. Explaining the reasons for this behavior is beyond the scope of this tutorial, and we refer the interested reader to von Luxburg, Belkin, and Bousquet (2004).

Which normalized Laplacian?

Looking at the differences between the two normalized spectral clustering algorithms using L_{rw} and L_{sym}, all three explanations of spectral clustering are in favor of L_{rw}. The reason is that the eigenvectors of L_{rw} are cluster indicator vectors $\mathbb{1}_{A_i}$, while the eigenvectors of L_{sym} are additionally multiplied with $D^{1/2}$, which might lead to undesired artifacts. As using L_{sym} also does not have any computational advantages, we thus advocate for using L_{rw}.

9   Outlook and further reading

Spectral clustering goes back to Donath and Hoffman (1973), who first suggested to construct graph partitions based on eigenvectors of the adjacency matrix. In the same year, Fiedler (1973) discovered that bi-partitions of a graph are closely connected with the second eigenvector of the graph Laplacian, and he suggested to use this eigenvector to partition a graph. Since then, spectral clustering has been discovered, re-discovered, and extended many times in different communities, see for example Pothen, Simon, and Liou (1990), Simon (1991), Bolla (1991), Hagen and Kahng (1992), Hendrickson and Leland (1995), Van Driessche and Roose (1995), Barnard, Pothen, and Simon (1995), Spielman and Teng (1996), Guattery and Miller (1998). A nice overview over the

history of spectral clustering can be found in Spielman and Teng (1996).

In the machine learning community, spectral clustering became popular by the works of Shi and Malik (2000), Ng et al. (2002), Meila and Shi (2001), and Ding (2004). A huge number of papers has subsequently been published, dealing with various extensions, new applications, and theoretical results on spectral clustering. As examples consider spectral clustering applied to the co-clustering problem (Dhillon, 2001), spectral clustering with additional side information (Joachims, 2003) connections between spectral clustering and the weighted k-means algorithm (Dhillon, Guan, and Kulis, 2005), learning similarity functions based on spectral clustering (Bach and Jordan, 2004), an average case analysis of spectral clustering (McSherry, 2001), or spectral clustering in a distributed environment (Kempe and McSherry, 2004).

The success of spectral clustering is mainly based on the fact that it does not make any assumptions on the form of the clusters. As opposed to k-means, where the resulting clusters are always convex sets, spectral clustering can solve very general problems like intertwined spirals. Moreover, spectral clustering can be implemented efficiently even for large data sets, as long as we make sure that the similarity graph is sparse. Once the similarity graph is chosen, we just have to solve a linear problem, and there are no issues of getting stuck in local minima or restarting the algorithms for several times with different initializations. However, we have already mentioned that choosing a good similarity graph is not trivial, and spectral clustering can be quite unstable under different choices of the parameters for the neighborhood graphs. So spectral clustering cannot serve as a "black box algorithm" which automatically detects the correct clusters in any given data set. But it can be considered as a powerful tool which can produce extremely good results if applied with care.

In the field of machine learning, graph Laplacians are not only used for clustering, but also emerge for many other tasks such as semi-supervised learning (e.g., Chapelle, Schölkopf, and Zien, 2006 for an overview) or manifold reconstruction (e.g., Belkin and Niyogi, 2003). In most applications, graph Laplacians are used to encode the assumption that data points which are "close" (i.e., w_{ij} is large) should have a "similar" label (i.e., $f_i \approx f_j$). The function f satisfies this assumption if $w_{ij}(f_i - f_j)^2$ is small for all i, j, that is $f'Lf$ is small. With this intuition one can for example use the quadratic form $f'Lf$ as a regularizer in a transductive classification problem. One other way to interpret the use of graph Laplacians is by the smoothness assumptions they encode. A function f which has a low value of $f'Lf$ has the property that it varies only "a little bit" in regions where the data points lie dense (i.e., the graph is tightly connected), whereas it is allowed to vary more (e.g., to change the sign) in regions of low data density. In this sense, a small value of $f'Lf$ encodes the so called "cluster assumption" in semi-supervised learning, which requests that the decision boundary of a classifier should lie in a region of low density.

An intuition often used is that graph Laplacians formally look like a continuous Laplace operator (and this is also where the name "graph Laplacian" comes from). To see this, transform a local similarity w_{ij} to a distance d_{ij} by the relationship $w_{ij} = 1/d_{ij}^2$ and observe that

$$ w_{ij}(f_i - f_j)^2 \approx \left(\frac{f_i - f_j}{d_{ij}} \right)^2 $$

looks like a difference quotient. As a consequence, the equation $f'Lf = \sum_{ij} w_{ij}(f_i - f_j)^2$ from Proposition 1 looks like a discrete version of the quadratic form associated to the standard Laplace operator $\mathcal{L}$ on $\mathbb{R}^n$, which satisfies

$$ \langle g, \mathcal{L}g \rangle = \int |\nabla g|^2 dx. $$

This intuition has been made precise in the works of Belkin (2003), Lafon (2004), Hein, Audibert, and von Luxburg (2005), Belkin and Niyogi (2005), Hein (2006), Giné and Koltchinskii (2005). In general, it is proved that graph Laplacians are discrete versions of certain continuous Laplace operators, and that if the graph Laplacian is constructed on a similarity graph of randomly sampled data points, then it converges to some continuous Laplace operator (or Laplace-Beltrami operator) on the underlying space. Belkin (2003) studied the first step of the convergence proof, which deals with the convergence of a continuous operator related to discrete graph Laplacians to the Laplace-Beltrami operator.His results were generalized from uniform distributions to general

distributions by Lafon (2004). Then in Belkin and Niyogi (2005), the authors prove pointwise convergence results for the unnormalized graph Laplacian using the Gaussian similarity function on manifolds with uniform distribution. At the same time, Hein et al. (2005) prove more general results, taking into account all different graph Laplacians L, L_{rw}, and L_{sym}, more general similarity functions, and manifolds with arbitrary distributions. In Giné and Koltchinskii (2005), distributional and uniform convergence results are proved on manifolds with uniform distribution. Hein (2006) studies the convergence of the smoothness functional induced by the graph Laplacians and shows uniform convergence results.

Apart from applications of graph Laplacians to partitioning problems in the widest sense, graph Laplacians can also be used for completely different purposes, for example for graph drawing (Koren, 2005). In fact, there are many more tight connections between the topology and properties of graphs and the graph Laplacian matrices than we have mentioned in this tutorial. Now equipped with an understanding for the most basic properties, the interested reader is invited to further explore and enjoy the huge literature in this field on his own.

References

Aldous, D. and Fill, J. (in preparation). *Reversible Markov Chains and Random Walks on Graphs.* online version available at http://www.stat.berkeley.edu/users/aldous/RWG/book.html.

Bach, F. and Jordan, M. (2004). Learning spectral clustering. In S. Thrun, L. Saul, and B. Schölkopf (Eds.), *Advances in Neural Information Processing Systems 16 (NIPS).* Cambridge, MA: MIT Press.

Bapat, R., Gutman, I., and Xiao, W. (2003). A simple method for computing resistance distance. *Z. Naturforsch.*, *58*, 494 - 498.

Barnard, S., Pothen, A., and Simon, H. (1995). A spectral algorithm for envelope reduction of sparse matrices. *Numerical Linear Algebra with Applications*, *2*(4), 317–334.

Belkin, M. (2003). *Problems of Learning on Manifolds*. PhD Thesis, University of Chicago.

Belkin, M. and Niyogi, P. (2003). Laplacian eigenmaps for dimensionality reduction and data representation. *Neural Computation*, *15*(6), 1373-1396.

Belkin, M. and Niyogi, P. (2005). Towards a theoretical foundation for Laplacian-based manifold methods. In P. Auer and R. Meir (Eds.), *Proceedings of the 18th Annual Conference on Learning Theory (COLT).* Springer, New York.

Ben-David, S., von Luxburg, U., and Pal, D. (2006). A sober look on clustering stability. In G. Lugosi and H. Simon (Eds.), *Proceedings of the 19th Annual Conference on Learning Theory (COLT)* (p. 5 - 19). Springer, Berlin.

Ben-Hur, A., Elisseeff, A., and Guyon, I. (2002). A stability based method for discovering structure in clustered data. In *Pacific Symposium on Biocomputing.*

Bhatia, R. (1997). *Matrix Analysis*. Springer, New York.

Bie, T. D. and Cristianini, N. (2006). Fast SDP relaxations of graph cut clustering, transduction, and other combinatorial problems . *JMLR*, *7*, 1409–1436.

Bolla, M. (1991). *Relations between spectral and classification properties of multigraphs* (Technical Report No. DIMACS-91-27). Center for Discrete Mathematics and Theoretical Computer Science.

Brémaud, P. (1999). *Markov chains: Gibbs fields, Monte Carlo simulation, and queues*. New York: Springer-Verlag.

Bui, T. N. and Jones, C. (1992). Finding good approximate vertex and edge partitions is NP-hard. *Inf. Process. Lett.*, *42*(3), 153–159.

Chapelle, O., Schölkopf, B., and Zien, A. (Eds.). (2006). *Semi-Supervised Learning*. MIT Press, Cambridge.

Chung, F. (1997). *Spectral graph theory*. Washington: Conference Board of the Mathematical Sciences.

Dhillon, I. (2001). Co-clustering documents and words using bipartite spectral graph partitioning. In *Proceedings of the seventh ACM SIGKDD international conference on Knowledge discovery and data mining (KDD)* (pp. 269–274). New York: ACM Press.

Dhillon, I., Guan, Y., and Kulis, B. (2005). *A unified view of kernel k-means, spectral clustering, and graph partitioning* (Technical Report No. UTCS TR-04-25). University of Texas at Austin.

Ding, C. (2004). *A tutorial on spectral clustering.* Talk presented at ICML. (Slides available at `http://crd.lbl.gov/~cding/Spectral/`)

Ding, C., He, X., Zha, H., Gu, M., and Simon, H. (2001). A min-max cut algorithm for graph partitioning and data clustering. In *Proceedings of the first IEEE International Conference on Data Mining (ICDM)* (pp. 107–114). Washington, DC, USA: IEEE Computer Society.

Donath, W. E. and Hoffman, A. J. (1973). Lower bounds for the partitioning of graphs. *IBM J. Res. Develop.*, *17*, 420–425.

Fiedler, M. (1973). Algebraic connectivity of graphs. *Czechoslovak Math. J.*, *23*, 298–305.

Fouss, F., Pirotte, A., Renders, J.-M., and Saerens, M. (2006). *A novel way of computing dissimilarities between nodes of a graph, with application to collaborative filtering and subspace projection of the graph nodes.* Technical Report.

Giné, E. and Koltchinskii, V. (2005). Empirical graph Laplacian approximation of Laplace-Beltrami operators: large sample results. In *Proceedings of the 4th International Conference on High Dimensional Probability.*

Golub, G. and Van Loan, C. (1996). *Matrix computations*. Baltimore: Johns Hopkins University Press.

Guattery, S. and Miller, G. L. (1998). On the quality of spectral separators. *SIAM Journal of Matrix Anal. Appl.*, *19*(3), 701–719.

Gutman, I. and Xiao, W. (2004). Generalized inverse of the Laplacian matrix and some applications. *Bulletin de l'Academie Serbe des Sciences at des Arts (Cl. Math. Natur.)*, *129*, 15-23.

Hagen, L. and Kahng, A. (1992). New spectral methods for ratio cut partitioning and clustering. *IEEE Trans. Computer-Aided Design*, *11*(9), 1074-1085.

Hastie, T., Tibshirani, R., and Friedman, J. (2001). *The elements of statistical learning*. New York: Springer.

Hein, M. (2006). Uniform convergence of adaptive graph-based regularization. In *Proceedings of the 19th Annual Conference on Learning Theory (COLT)* (p. 50-64). Springer, New York.

Hein, M., Audibert, J.-Y., and von Luxburg, U. (2005). From graphs to manifolds - weak and strong pointwise consistency of graph Laplacians. In P. Auer and R. Meir (Eds.), *Proceedings of the 18th Annual Conference on Learning Theory (COLT)* (p. 470-485). Springer, New York.

Hendrickson, B. and Leland, R. (1995). An improved spectral graph partitioning algorithm for mapping parallel computations. *SIAM J. on Scientific Computing*, *16*, 452–469.

Joachims, T. (2003). Transductive Learning via Spectral Graph Partitioning. In T. Fawcett and N. Mishra (Eds.), *Proceedings of the 20th international conference on machine learning (ICML)* (p. 290-297). AAAI Press.

Kempe, D. and McSherry, F. (2004). A decentralized algorithm for spectral analysis. In *Proceedings of the 36th Annual ACM Symposium on Theory of Computing (STOC)* (pp. 561–568). New York, NY, USA: ACM Press.

Klein, D. and Randic, M. (1993). Resistance distance. *Journal of Mathematical Chemistry*, *12*, 81-95.

Koren, Y. (2005). Drawing graphs by eigenvectors: theory and practice. *Computers and Mathematics with Applications*, *49*, 1867–1888.

Lafon, S. (2004). *Diffusion maps and geometric harmonics*. PhD Thesis, Yale University.

Lang, K. (2006). Fixing two weaknesses of the spectral method. In Y. Weiss, B. Schölkopf, and J. Platt (Eds.), *Advances in Neural Information Processing Systems 18* (pp. 715–722). Cambridge, MA: MIT Press.

Lange, T., Roth, V., Braun, M., and Buhmann, J. (2004). Stability-based validation of clustering solutions. *Neural Computation*, *16*(6), 1299 – 1323.

Lovász, L. (1993). Random walks on graphs: a survey. In *Combinatorics, Paul Erdos is eighty, Vol. 2 (Keszthely, 1993)* (Vol. 2, pp. 353–397). Budapest: János Bolyai Math. Soc.

Lütkepohl, H. (1997). *Handbook of Matrices*. Chichester: Wiley.

McSherry, F. (2001). Spectral partitioning of random graphs. In *Proceedings of the 42nd IEEE symposium on Foundations of Computer Science (FOCS)* (p. 529). Washington, DC, USA: IEEE Computer Society.

Meila, M. and Shi, J. (2001). A random walks view of spectral segmentation. In *8th International Workshop on Artificial Intelligence and Statistics (AISTATS).*

Mohar, B. (1991). The Laplacian spectrum of graphs. In *Graph theory, combinatorics, and applications. Vol. 2 (Kalamazoo, MI, 1988)* (pp. 871–898). New York: Wiley.

Mohar, B. (1997). Some applications of Laplace eigenvalues of graphs. In G. Hahn and G. Sabidussi (Eds.), *Graph Symmetry: Algebraic Methods and Applications* (Vol. NATO ASI Ser. C 497, p. 225-275). Kluwer.

Nadler, B., Lafon, S., Coifman, R., and Kevrekidis, I. (2006). Diffusion maps, spectral clustering and eigenfunctions of Fokker-Planck operators. In Y. Weiss, B. Schölkopf, and J. Platt (Eds.), *Advances in Neural Information Processing Systems 18* (pp. 955–962). Cambridge, MA: MIT Press.

Ng, A., Jordan, M., and Weiss, Y. (2002). On spectral clustering: analysis and an algorithm. In T. Dietterich, S. Becker, and Z. Ghahramani (Eds.), *Advances in Neural Information Processing Systems 14.* MIT Press.

Norris, J. (1997). *Markov Chains*. Cambridge: Cambridge University Press.

Pothen, A., Simon, H. D., and Liou, K. P. (1990). Partitioning sparse matrices with eigenvectors of graphs. *SIAM Journal of Matrix Anal. Appl.*, *11*, 430–452.

Shi, J. and Malik, J. (2000). Normalized cuts and image segmentation. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, *22*(8), 888-905.

Simon, H. (1991). Partitioning of unstructured problems for parallel processing. *Computing Systems Engineering*, *2*, 135-148.

Spielman, D. and Teng, S. (1996). Spectral partitioning works: planar graphs and finite element meshes. In *37th Annual Symposium on Foundations of Computer Science (Burlington, VT, 1996)* (pp. 96–105). Los Alamitos, CA: IEEE Comput. Soc. Press.

Stewart, G. and Sun, J. (1990). *Matrix Perturbation Theory*. New York: Academic Press.

Stoer, M. and Wagner, F. (1997). A simple min-cut algorithm. *J. ACM*, *44*(4), 585–591.

Tibshirani, R., Walther, G., and Hastie, T. (2001). Estimating the number of clusters in a dataset via the gap statistic. *J. Royal. Statist. Soc. B*, *63*(2), 411–423.

Van Driessche, R. and Roose, D. (1995). An improved spectral bisection algorithm and its application to dynamic load balancing. *Parallel Comput.*, *21*(1), 29–48.

von Luxburg, U., Belkin, M., and Bousquet, O. (2004). *Consistency of spectral clustering* (Technical Report No. 134). Max Planck Institute for Biological Cybernetics.

von Luxburg, U., Bousquet, O., and Belkin, M. (2004). On the convergence of spectral clustering on random samples: the normalized case. In J. Shawe-Taylor and Y. Singer (Eds.), *Proceedings of the 17th Annual Conference on Learning Theory (COLT)* (p. 457-471). Springer, New York.

von Luxburg, U., Bousquet, O., and Belkin, M. (2005). Limits of spectral clustering. In L. Saul, Y. Weiss, and L. Bottou (Eds.), *Advances in Neural Information Processing Systems (NIPS) 17* (p. 857-864). Cambridge, MA: MIT Press.

Wagner, D. and Wagner, F. (1993). Between min cut and graph bisection. In *Proceedings of the 18th International Symposium on Mathematical Foundations of Computer Science (MFCS)* (pp. 744–750). London: Springer.